



pg_index_stats

manage extended statistics automatically

Lepikhov A., Rybakina A.

Postgres Professional

2024

Self Introduction

- Core Developer in Postgres Professional since 2017.
- Contributing to the PostgreSQL project since 2017:
Self-Join Elimination, GROUP-BY optimisation, OR <-> ANY transformation.
- Ph.D. in Computer Sciences (Distributed Databases), MSU, 2008.
- Designed extensions: AQO, sr_plan, pg_index_stats ...



Extended Statistics

CREATE STATISTICS ON x1,x2,x3,x4 FROM tablename;

- **MCV** - Most Common Values on composite value of (x1,x2,x3).
- **ndistinct** - number of distinct values on all combinations of columns: (x1,x2),(x1,x3),(x1,x2,x3)...
- **dependencies** - functional dependencies between combinations of columns: $x1 \rightarrow x2$, $(x1,x2) \rightarrow x3$,...



Vondra. T. CREATE STATISTICS improvements,
PGConf.DE 2022

PostgresPro

Laboriousness of Extended Statistics

- **MCV** - two arrays: `values[]` and `frequencies[]`.
- **ndistinct** - 1 integer for each of $2^n - (n + 1)$ combinations
- **dependencies** - 1 float value for each of combinations

columns:	2	3	4	...	8
distinct combinations:	1	4	11	...	247
dependency combinations:	2	9	28	...	1016

Quick start

```
CREATE EXTENSION 'pg_index_stats';
CREATE TABLE test (x int,y int,z text);
CREATE INDEX ON test (x,y);
CREATE INDEX ON test (z,y);
\dX
```

List of extended statistics						
Schema	Name	Definition	Ndistinct	Dependencies	MCV	
public	test_x_y_stat	x, y FROM test	defined	defined	defined	
public	test_z_y_stat	y, z FROM test	defined	defined	defined	



Quick start - II

```
DROP INDEX test_x_y_idx;
```

List of extended statistics						
Schema	Name	Definition	Ndistinct	Dependencies	MCV	
public	test_z_y_stat	y, z FROM test	defined	defined	defined	

```
SELECT stxname,obj_description(oid,'pg_statistic_ext')
FROM (SELECT oid,stxname FROM pg_statistic_ext);
```

stxname	obj_description
test_z_y_stat	pg_index_stats - multivariate statistics

GUCs & Funcs

- *mode* - disabled | all | univariate | multivariate
- *columns_limit*
- *pg_index_stats_build(relname, mode)*
- *pg_index_stats_rebuild()*
- *pg_index_stats_remove()*

Two-step process:

- Object Access Hook - gather candidate OIDs
- Utility Hook - create extended statistics after a successful utility statement

Set dependency of auto-generated statistics on the extension and index relation.

Add specific description for auto-generated statistics

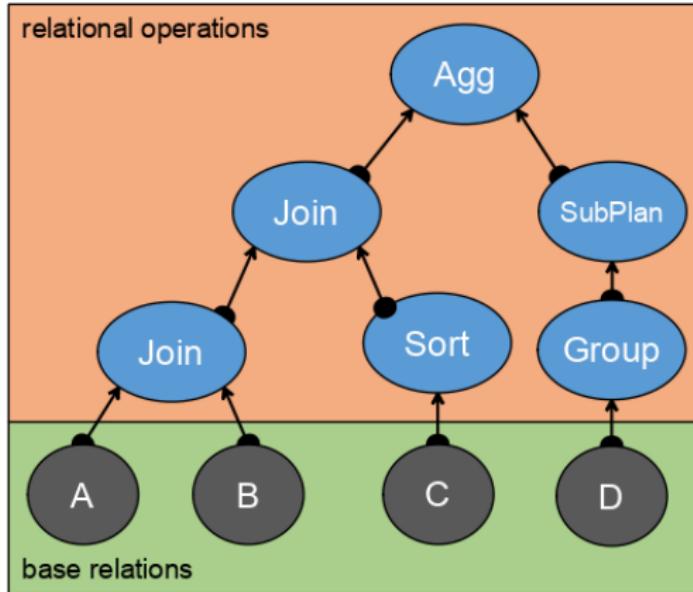
Query auto-generation frameworks



Serve a RESTful API from any Postgres database



What is the reason?



Cardinality estimation:

- Adaptive Query Optimizer
(*aqo*)
 - Query replanning
(*replan*)
 - Plan freezing
(*sr_plan*)
 - Selectivity by index
(*JoinSel*)

Cardinality estimation:

- ## ➤ statistics ???

Redundant expressions

Query	Planned rows	Actual rows
SELECT * FROM power_plants WHERE country = 'RUS';	544	544
SELECT * FROM power_plants WHERE country = 'RUS' AND country_long = 'Russia';	8	544



*Global Power Plant Database
Copyright 2018-2021 World Resources Institute and Data Contributors

The question

Why can database systems, which manage all the data, not analyse it and find at least in-table dependencies and interconnections?



Why indexes?

Table "Parcels":

Indexes:

- "parcel_pkey" **PRIMARY KEY**, btree (id)
- "parcel_parcel_id" btree (parcel_id)
- "parcel_id_prik" btree (parcel_id, prik)
- "parcel_id_sn_pol" btree (parcel_id, sn_pol)
- "parcel_par_begin" btree (parcel_id, prik_begin_period)
- "parcel_par_patient" btree (parcel_id, patient_id)
- "parcel_par_recid" btree (parcel_id, recid)
- "parcel_per_prik" btree (period, prik)

Extended v/s Plain Statistics

SELECT * FROM power_plants WHERE ...

Query	Plain stat	Extended stat	Actual rows
country = 'RUS';	544	545	544
AND primary_fuel IN ('Solar')	166	57	57
primary_fuel IN ('Solar', 'Biomass')	189	60	60
primary_fuel IN ('Solar', 'Biomass', 'Coal')	225	156	156
AND source = 'Wiki—Solar'	24	46	40
AND longitude > 40.	1	1	33
AND longitude < 70.			



*Global Power Plant Database

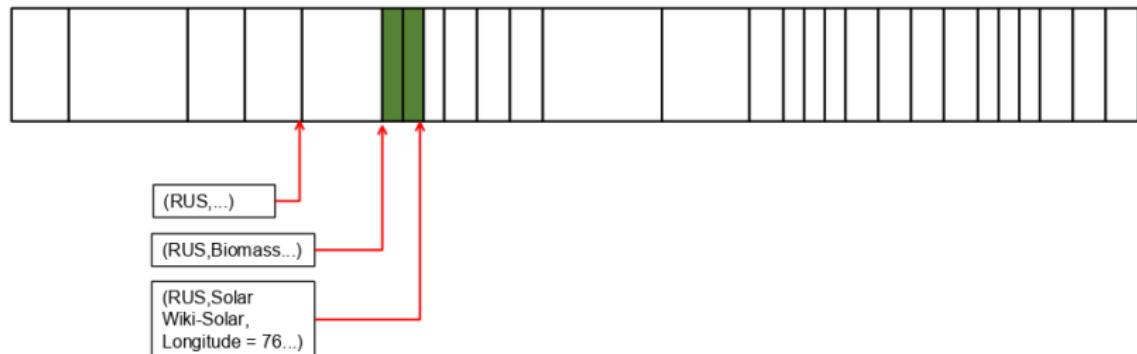
Copyright 2018-2021 World Resources Institute and Data Contributors



Does multicolumn histogram make sense?

In general - no. We need a multidimensional histogram. But it can make sense if we follow the definition of indexes ...

```
SELECT * FROM power_plants
WHERE
    country = 'RUS' AND
    primary_fuel IN ('Solar', 'Biomass', 'Coal') AND
    AND source = 'Wiki-Solar' AND
    longitude > 40. AND longitude < 70.;
```



Multicolumn histogram advantage

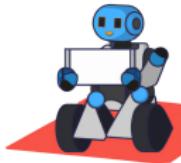
SELECT * FROM power_plants WHERE ...

Query	Plain stat	Extended stat	Histogram	Actual rows
country = 'RUS';	544	545	544	544
AND primary_fuel IN ('Solar')	166	57	57	57
primary_fuel IN ('Solar', 'Biomass')	189	60	60	60
primary_fuel IN ('Solar', 'Biomass', 'Coal')	225	156	156	156
AND source = 'Wiki-Solar'	24	46	42	40
AND longitude > 40.	1	1	35	33
AND longitude < 70.				



*Global Power Plant Database
Copyright 2018-2021 World Resources Institute and Data Contributors





Questions ?



The `pg_index_stats` extension
github link



Postgres Professional LLC
The Russian Postgres Company