



Query plan freezing extension design, issues and lessons learned

Belyalov D., Lepikhov A., Rybakina A.

Postgres Professional

2023

Self Introduction

- Specialist in Applied Mathematics graduated from Chelyabinsk State University in 2005
- Ph.D. in Computer Sciences (Distributed Databases) awarded at Moscow State University in 2008.
- Working in Postgres Professional as a Core Developer since 2017.
- Designed the Shardman project architecture on its earlier steps
- Worked on Multimaster project
- Working on various query optimization issues

Who we are

- Research team, part of Postgres Professional, dealing with optimization issues
- Caused by the idea of sustainable coding
- Design enterprise and core features to improve the planner effectiveness
- Projects: Self-Join Removal, Asymmetric JOIN, Optimized Group-by, AQO, sr_plan ...

Plan freezing?

Reason 1:

Don't optimize next time!

pgbench:

- With planning: ≈ 6000 tps
- Prepared statements: ≈ 7600 tps



- Stale statistics
- Imperfection of cost estimation algorithms
- Implicit functional dependencies between columns

Functional Dependencies (overestimation)

```
CREATE TABLE people (  
  name          text ,  
  occupation    text ,  
  sex           boolean ,  
  region        text ,  
  is_vaccinated boolean  
);
```

```
SELECT * FROM people t1  
WHERE occupation = 'Tractor Driver' AND sex = 'female'  
/*  
Seq Scan on people (rows=1584) (actual rows=1)  
  Filter: (is_woman AND (occupation = 'Tractor Driver'))  
  Rows Removed by Filter: 99999  
*/
```

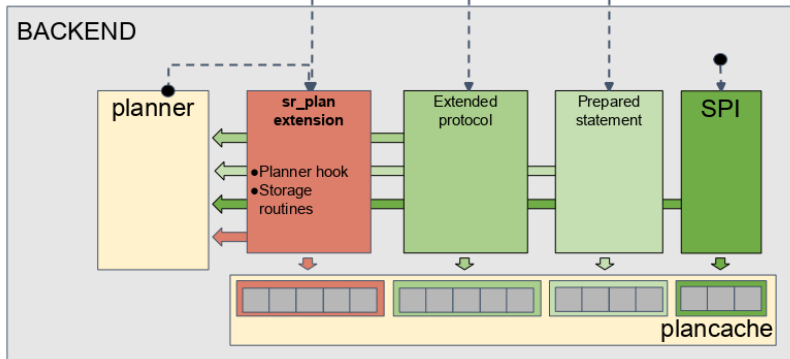
Functional Dependencies (underestimation)

```
CREATE TABLE people (  
  name          text ,  
  occupation    text ,  
  sex           boolean ,  
  region        text ,  
  is_vaccinated boolean  
);
```

```
SELECT * FROM people  
WHERE region = 'Chelyabinsk' AND is_vaccinated;  
/*  
Seq Scan on people (rows=114) (actual rows=907)  
  Filter: (is_vaccinated AND (region = 'Chelyabinsk'))  
  Rows Removed by Filter: 99094  
*/
```


Unfortunate planning is inevitable [at least, for now]. Assuming someone or something could force the planner to generate a better plan, we should provide the tool to freeze the right solution for the subsequent executions.

Stick on the plan in the Plan Cache ?



The sr_plan extension

- Abbreviates **save/restore** plan
- Introduced in Postgres Pro Enterprise 15 (don't mix up with the extension sr_plan existed up to PGPro Enterprise 13!)
- Freezes specific plan for a [parameterized] query



How it works

- `sr_register_query('SELECT ... WHERE x = $1 AND y = 42', ...)`
- `sr_plan_freeze(srid)`
- `sr_plan_unfreeze(srid)`

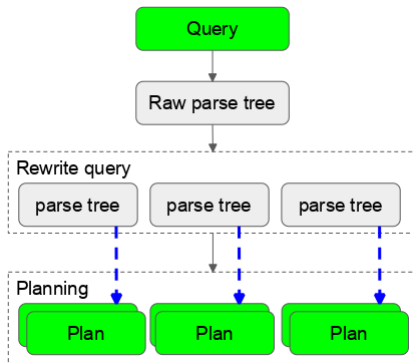
Frozen plan have a special node in explain

QUERY PLAN

```
-----  
Custom Scan (SRScan) (actual rows=1 loops=1)  
  Frozen plan ID: 1  
    -> Aggregate (actual rows=1 loops=1)  
      -> Seq Scan on a (actual rows=10 loops=1)  
        Filter: ((x = \ $1) AND (y = 42))
```


Lesson 1

In DBMS, the way from a query text to the plan is not straightforward.




Lesson 2

Only one way to prove applicability of the plan to the given query is to compare stored and incoming parse trees



Lesson 3


To make overhead admissible,
we should have kind of parse
tree signature - queryId



Lesson 4


To apply plan freezing for parameterized queries we should '*generalize*' the parse tree

```
sr_register_query('SELECT ... FROM A, ... WHERE x = $1...')
```




Lesson 5

The extension loading
order matters



Points of overhead

- Parse tree comparison
- Plan invalidation
 - Per-backend cache invalidation
 - Disc storage sync
 - Transactional issues



Freezing procedure


```
SELECT srid FROM sr_register_query('
  SELECT abalance FROM pgbench_accounts
  WHERE aid = $1', 'int') \gset
SELECT abalance FROM pgbench_accounts WHERE aid = 1;
SELECT sr_plan_freeze(:srid);
```

```
UPDATE pgbench_accounts ...
UPDATE pgbench_tellers ...
UPDATE pgbench_branches ...
INSERT INTO pgbench_history ...
```

- With planning: ≈ 6000 tps
- Frozen statements: ≈ 6500 tps

The future

- Detect bad plan and try something different
- Plan transfer procedure
- Global prepared statements



Questions ?